

1 Numerical Tests

We implement the Metropolis algorithm in R. Inside the Metropolis algorithm, we evaluate the likelihood function using the DTQ method, which is implemented in C++ as an R package. Note that all code and data used in this work is available online (<https://github.com/hbhat4000/sdeinference>)—see the “Rdtq2d” and “pursuit2d” directories. To test the method, we first consider the SDE

$$dX_{1,t} = -\frac{X_{2,t}}{L}dt + \frac{s_1^2}{L}dW_{1,t}, \quad dX_{2,t} = \frac{X_{1,t}}{C}dt + \frac{s_2^2}{C}dW_{2,t}. \quad (1)$$

This system describes a noisy electrical oscillator with one inductor (with inductance L) and one capacitor (with capacitance C). The dependent variables $X_{1,t}$ and $X_{2,t}$ represent, respectively, the current and voltage of the circuit at time t .

Our goal here is to test the performance of the algorithm using simulated data. To generate this data, we start with known values of the parameters: $L = C = (2\pi)^{-1}$ and $s_1 = s_2 = .4/\sqrt{2\pi}$. Using a fixed initial condition $(X_{1,0}, X_{2,0})$, we then use the Euler-Maruyama method to step (??) forward in time until a final time $T > 0$. When we carry out this time-stepping, we use a step size of 0.001 and then retain only those samples at times $t_m = m\Delta t$, from $m = 0$ to $m = M$, where $M\Delta t = T$. The simulated data is taken over two periods of the oscillator ($T = 2$) with a full resolution of $\Delta t = 0.01$. By, for example, taking every other row of this data set, we can obtain data with a resolution of $\Delta t = 0.02$.

Using the samples $\{\mathbf{x}_m\}_{m=0}^M$ thus constructed, we run the Metropolis algorithm. Because capacitance and inductance are physically constrained to be positive, we set $1/L = \theta_1^2$. For the tests presented here, we infer only θ_1 , keeping other parameters fixed at their known values. For θ_1 , we use a diffuse Gaussian prior with mean 0 and standard deviation 100. For the proposal distribution Z_{N+1} in the auxiliary Markov chain, we choose i.i.d. Gaussians with mean 0 and standard deviation 0.35.

When we run the Metropolis algorithm, we discard the first 100 samples and retain the next 1000 samples. For each value of Δt and the DTQ time step h , we compute both the mean of the samples of θ_1^2 and the mode of the kernel density estimate of θ_1^2 . We compare these values against the true value of the parameter $1/L = 2\pi$ and record the relative errors as, respectively, e_1 and e_2 :

Δt	h	e_1 (relative error of mean)	e_2 (relative error of mode)
0.04	0.04	6.1%	7.6%
0.04	0.02	0.54%	6.8%
0.04	0.01	5.1%	1.1%
0.02	0.02	12%	14%
0.02	0.01	4.9%	2.3%

When $h = \Delta t$, only one step of the method described in Section ?? is required to go from time t_m to t_{m+1} . This step does not use any quadrature at all—one merely evaluates (??) using a point mass for the density at time t_m . The resulting likelihood function is a product of Gaussians. On the other hand, when h is strictly less than Δt , we must use quadrature (i.e., the actual DTQ method) to step forward in time

from t_m to t_{m+1} . Clearly, using the DTQ method to compute the likelihood yields more accurate posteriors than using a purely Gaussian likelihood.

To visualize the results, we present Figure ?? . The true value of $1/L = 2\pi$ is indicated by the dashed black line. The posterior mode for $h = 0.01$ is indicated by the solid black line. The curves are kernel density estimates computed using the posterior samples described above.




Fig. 1 We plot kernel density estimates of posterior densities $p(\theta_1^2 | \mathbf{x})$. We use simulated data with $\Delta t = 0.04$, generated as described above. Each posterior density corresponds to a finer DTQ step h . As we take a finer DTQ step (i.e., as h decreases), the posterior mode approaches the true value indicated by the solid vertical line at $1/L = 2\pi$.